


MoE Training on Hopper GPUs



High activation
memory

Communication
overhead

MXFP4-Based Training Techniques



Efficiency Gains

